

Lab 5: Creating a Model Service

In this lab, we'll use the churn prediction model trained on a smaller set of features.

Question 1

Install Pipenv

What's the version of pipenv you installed?

Use --version to find out

Question 2

Use Pipenv to install Scikit-Learn version 1.0

What's the first hash for scikit-learn you get in Pipfile.lock?

Models

We've prepared a dictionary vectorizer and a model.

They were trained (roughly) using this code:

```
features = ['tenure', 'monthlycharges', 'contract']  
dicts = df[features].to_dict(orient='records')
```

```
dv = DictVectorizer(sparse=False)  
X = dv.fit_transform(dicts)
```

```
model = LogisticRegression().fit(X, y)
```

Note: You don't need to train the model. This code is just for your reference.

And then saved with Pickle. Download them:

DictVectorizer

LogisticRegression

With wget:

```
PREFIX=https://raw.githubusercontent.com/fenago/mlbookcamp-code/master/course-  
zoomcamp/05-deployment/lab  
wget $PREFIX/model1.bin  
wget $PREFIX/dv.bin
```

Question 3

Let's use these models!

Write a script for loading these models with pickle

Score this customer:

```
{"contract": "two_year", "tenure": 12, "monthlycharges": 19.7}
```

What's the probability that this customer is churning?

If you're getting errors when unpickling the files, check their checksum:

```
$ md5sum model1.bin dv.bin
5868e129bfbb309ba60bf750263afab1 model1.bin
c49b69f8a5a3c560882ff5daa3c0ff4d dv.bin
```

Question 4

Now let's serve this model as a web service

Install Flask and Gunicorn (or waitress, if you're on Windows)

Write Flask code for serving the model

Now score this customer using requests:

```
url = "YOUR_URL"
```

```
customer = {"contract": "two_year", "tenure": 1, "monthlycharges": 10}
```

```
requests.post(url, json=customer).json()
```

What's the probability that this customer is churning?

Docker

Install Docker. We will use it for the next two questions.

For these questions, I prepared a base image: `agrigorev/zoomcamp-model:3.8.12-slim`. You'll need to use it (see Question 5 for an example).

This image is based on `python:3.8.12-slim` and has a logistic regression model (a different one) as well a dictionary vectorizer inside.

This is how the Dockerfile for this image looks like:

```
FROM python:3.8.12-slim
```

```
WORKDIR /app
```

```
COPY ["model2.bin", "dv.bin", "./"]
```

I already built it and then pushed it to `agrigorev/zoomcamp-model:3.8.12-slim`.

Note: You don't need to build this docker image, it's just for your reference.

Question 5

Now create your own Dockerfile based on the image I prepared.

It should start like that:

```
FROM agrigorev/zoomcamp-model:3.8.12-slim
```

```
# add your stuff here
```

Now complete it:

Install all the dependencies from the Pipenv file

Copy your Flask script

Run it with gunicorn

When you build your image, what's the image id for agrigorev/zoomcamp-model:3.8.12-slim?

Look at the first step of your build log. It should look something like that:

```
$ docker some-command-for-building
```

```
Sending build context to Docker daemon 2.048kB
```

```
Step 1/N : FROM agrigorev/zoomcamp-model:3.8.12-slim
```

```
---> XXXXXXXXXXXXXXX
```

```
Step 2/N : ....
```

```
You need this XXXXXXXXXXXXXXX.
```

Alternatively, you can get this information when running docker images - it'll be in the "IMAGE ID" column. Submitting DIGEST (long string starting with "sha256") is also fine.

Question 6

Let's run your docker container!

After running it, score this customer:

```
url = "YOUR_URL"
```

```
customer = {"contract": "two_year", "tenure": 12, "monthlycharges": 10}
```

```
requests.post(url, json=customer).json()
```

What's the probability that this customer is churning?