

# Lab 7: Neural Networks

## Dataset

In this lab, we'll build a model for predicting if we have an image of a dog or a cat. For this, we will use the "Dogs & Cats" dataset that can be downloaded from [Kaggle](https://www.kaggle.com/c/dogs-vs-cats/data).

You need to download the `train.zip` file.

If you have trouble downloading from Kaggle, use [this link](https://github.com/fenago/large-datasets/releases/download/dogs-cats/train.zip) instead:

```
```bash
wget https://github.com/fenago/large-datasets/releases/download/dogs-cats/train.zip
```
```

In the lectures we saw how to use a pre-trained neural network. In the lab, we'll train a much smaller model from scratch.

**\*\*Note:\*\*** You don't need a computer with a GPU for this lab. A laptop or any personal computer should be sufficient.

## Data Preparation

The dataset contains 12,500 images of cats and 12,500 images of dogs.

Now we need to split this data into train and validation

- \* Create a `train` and `validation` folders
- \* In each folder, create `cats` and `dogs` folders
- \* Move the first 10,000 images to the train folder (from 0 to 9999) for both cats and dogs - and put them in respective folders
- \* Move the remaining 2,500 images to the validation folder (from 10000 to 12499)

You can do this manually or with Python (check `os` and `shutil` packages).

### ### Model

For this lab we will use Convolutional Neural Network (CNN. Like in the lectures, we'll use Keras.

You need to develop the model with following structure:

- \* The shape for input should be `(150, 150, 3)`
- \* Next, create a convolutional layer  
(`Conv2D`)([https://keras.io/api/layers/convolution\\_layers/convolution2d/](https://keras.io/api/layers/convolution_layers/convolution2d/)):
  - \* Use 32 filters
  - \* Kernel size should be `(3, 3)` (that's the size of the filter)
  - \* Use `relu` as activation
- \* Reduce the size of the feature map with max pooling  
(`MaxPooling2D`)([https://keras.io/api/layers/pooling\\_layers/max\\_pooling2d/](https://keras.io/api/layers/pooling_layers/max_pooling2d/))
  - \* Set the pooling size to `(2, 2)`
- \* Turn the multi-dimensional result into vectors using a  
(`Flatten`)([https://keras.io/api/layers/reshaping\\_layers/flatten/](https://keras.io/api/layers/reshaping_layers/flatten/)) layer
- \* Next, add a `Dense` layer with 64 neurons and `relu` activation
- \* Finally, create the `Dense` layer with 1 neuron - this will be the output
  - \* The output layer should have an activation - use the appropriate activation for the binary classification case

As optimizer use (`SGD`)(<https://keras.io/api/optimizers/sgd/>) with the following parameters:

- \* `SGD(lr=0.002, momentum=0.8)`

For clarification about kernel size and max pooling, check [Week #11 Office Hours](<https://www.youtube.com/watch?v=1WRgdBTUaAc>).

## Question 1

Since we have a binary classification problem, what is the best loss function for us?

Note: since we specify an activation for the output layer, we don't need to set `from\_logits=True`

## Question 2

What's the total number of parameters of the model? You can use the ``summary`` method for that.

### ### Generators and Training

For the next two questions, use the following data generator for both train and validation:

```
```python
ImageDataGenerator(rescale=1./255)
```
```

- \* We don't need to do any additional pre-processing for the images.
- \* When reading the data from train/val directories, check the ``class_mode`` parameter. Which value should it be for a binary classification problem?
- \* Use ``batch_size=20``
- \* Use ``shuffle=True`` for both training and validation

For training use ``fit()`` with the following params:

```
```python
model.fit(
    train_generator,
    steps_per_epoch=100,
    epochs=10,
    validation_data=validation_generator,
    validation_steps=50
)
```
```

Note ``validation_steps=50`` - this parameter says "run only 50 steps on the validation data for evaluating the results".

This way we iterate a bit faster, but don't use the entire validation dataset.

That's why it's important to shuffle the validation dataset as well.

## Question 3

What is the median of training accuracy for this model?

## Question 4

What is the standard deviation of training loss for this model?

### Data Augmentation

For the next two questions, we'll generate more data using data augmentations.

Add the following augmentations to your training data generator:

```
* `rotation_range=40`,`  
* `width_shift_range=0.2`,`  
* `height_shift_range=0.2`,`  
* `shear_range=0.2`,`  
* `zoom_range=0.2`,`  
* `horizontal_flip=True`,`  
* `fill_mode='nearest'`
```

## Question 5

Let's train our model for 10 more epochs using the same code as previously. Make sure you don't re-create the model - we want to continue training the model we already started training.

What is the mean of validation loss for the model trained with augmentations?

## Question 6

What's the average of validation accuracy for the last 5 epochs (from 6 to 10) for the model trained with augmentations?